

AHMAD ZAFAR AGAH

LAB1

10/8/24

[1.2 ARP, Wireshark, Netsim](#)

[1.2.1 ARP](#)

[1.2.2 ARP2](#)

[1.2.3 ARP \(Cloud\)](#)

[1.2.4 Netsim](#)

[1.3 Cloud Computing](#)

[1.3.3 Scan targets for services](#)

[1.3.5 Navigating default networks](#)

[1.3.6 Creating custom networks](#)

1.2 ARP, Wireshark, Netsim

1.2.1 ARP

```
agah@ada:~$ ip address
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host noprefixroute
        valid_lft forever preferred_lft forever
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 52:54:00:13:a0:c6 brd ff:ff:ff:ff:ff:ff
    altname enp0s3
    inet 131.252.208.103/24 metric 100 brd 131.252.208.255 scope global dynamic ens3
        valid_lft 12671sec preferred_lft 12671sec
agah@ada:~$
```

- What is the default router's IP address (e.g. the gateway address for the default route 0.0.0.0/0)

```
agah@ada:~$ netstat -rn
Kernel IP routing table
Destination        Gateway            Genmask           Flags   MSS Window  irtt Iface
0.0.0.0            131.252.208.1     0.0.0.0           UG        0 0          0 ens3
10.218.208.100     131.252.208.1     255.255.255.255  UGH        0 0          0 ens3
10.218.208.108     131.252.208.1     255.255.255.255  UGH        0 0          0 ens3
131.252.110.102    131.252.208.1     255.255.255.255  UGH        0 0          0 ens3
131.252.110.103    131.252.208.1     255.255.255.255  UGH        0 0          0 ens3
131.252.208.0      0.0.0.0           255.255.255.0     U          0 0          0 ens3
131.252.208.1      0.0.0.0           255.255.255.255  UH        0 0          0 ens3
131.252.208.53     0.0.0.0           255.255.255.255  UH        0 0          0 ens3
agah@ada:~$
```

- What is the name of the default router and its hardware address?

```
agah@ada:~$ arp 131.252.208.1
Address default router HWtype HWaddress hardware address Flags Mask Iface
router.seas.pdx.edu ether 00:00:5e:00:01:01 C
agah@ada:~$ arp -n 131.252.208.1
Address HWtype HWaddress Flags Mask Iface
131.252.208.1 ether 00:00:5e:00:01:01 C
agah@ada:~$
```

- How many entries are there in the ARP table?

```
agah@ada:~$ arp -a | wc -l
25
agah@ada:~$
```

1.2.2 ARP2

- List any IP addresses share the same hardware address

None of the IP addresses share the same hardware address (MAC)

```
agah@ada:~$ arp -a | sort -k 4
router.seas.pdx.edu (131.252.208.1) at 00:00:5e:00:01:01 [ether] on ens3
mirrors.cat.pdx.edu (131.252.208.20) at 00:00:5e:00:01:14 [ether] on ens3
cs162lab.cs.pdx.edu (131.252.208.81) at 00:00:5e:00:01:51 [ether] on ens3
cs302lab.cs.pdx.edu (131.252.208.83) at 00:00:5e:00:01:53 [ether] on ens3
cs163lab.cs.pdx.edu (131.252.208.84) at 00:00:5e:00:01:54 [ether] on ens3
gitlab.cecs.pdx.edu (131.252.208.138) at 00:00:5e:00:01:8a [ether] on ens3
jammy.cecs.pdx.edu (131.252.208.11) at 52:54:00:59:3e:39 [ether] on ens3
babbage.cs.pdx.edu (131.252.208.23) at 52:54:00:5c:6f:6e [ether] on ens3
focal.cecs.pdx.edu (131.252.208.94) at 52:54:00:78:73:00 [ether] on ens3
tanto.cs.pdx.edu (131.252.208.5) at 52:54:00:87:21:c4 [ether] on ens3
quizor6.cs.pdx.edu (131.252.208.60) at 52:54:00:a3:46:7f [ether] on ens3
dc-rdns-01.cat.pdx.edu (131.252.208.117) at 52:54:00:a9:30:9f [ether] on ens3
rdns.cat.pdx.edu (131.252.208.53) at 52:54:00:a9:30:9f [ether] on ens3
danimoth.cat.pdx.edu (131.252.208.34) at 52:54:00:b4:6e:05 [ether] on ens3
gitlab-01.cecs.pdx.edu (131.252.208.137) at 52:54:00:c2:05:63 [ether] on ens3
quizor4.cs.pdx.edu (131.252.208.36) at 52:54:00:cf:4c:1b [ether] on ens3
rita.cecs.pdx.edu (131.252.208.28) at 52:54:00:eb:9a:42 [ether] on ens3
ruby.cecs.pdx.edu (131.252.208.85) at 52:54:00:f2:09:bc [ether] on ens3
mircle.cat.pdx.edu (131.252.208.54) at 52:54:00:f6:f8:54 [ether] on ens3
quizor1.cs.pdx.edu (131.252.208.171) at cc:aa:77:07:f2:7a [ether] on ens3
silverfish.cat.pdx.edu (131.252.208.77) at cc:aa:77:0b:76:be [ether] on ens3
destiny.cat.pdx.edu (131.252.208.17) at cc:aa:77:50:b9:5d [ether] on ens3
expn.cat.pdx.edu (131.252.208.110) at cc:aa:77:5f:de:0e [ether] on ens3
stargate.cat.pdx.edu (131.252.208.43) at cc:aa:77:ed:72:3e [ether] on ens3
mirapo.cat.pdx.edu (131.252.208.63) at cc:aa:77:f1:d3:21 [ether] on ens3
agah@ada:~$ arp -a | sort -k 4 | awk '{print $4, $2}' | uniq -D -f 0
agah@ada:~$ arp -a | sort -k 4 | awk '{print $4, $2}' | uniq -D -f 0
agah@ada:~$
```

The output is null, which means no duplicate values in hardware addresses

The result of `arp -a | wc -l` and `arp -a | sort -k 4 | awk '{print $4, $2}' | uniq | wc -l` is 25, they are equal which means there are equal number of IP addresses and MAC addresses

```

agah@ada:~$ arp -a | sort -k 4 | awk '{print $4, $2}' | uniq
00:00:5e:00:01:01 (131.252.208.1)
00:00:5e:00:01:14 (131.252.208.20)
00:00:5e:00:01:51 (131.252.208.81)
00:00:5e:00:01:53 (131.252.208.83)
00:00:5e:00:01:54 (131.252.208.84)
00:00:5e:00:01:8a (131.252.208.138)
52:54:00:59:3e:39 (131.252.208.11)
52:54:00:5c:6f:6e (131.252.208.23)
52:54:00:78:73:00 (131.252.208.94)
52:54:00:87:21:c4 (131.252.208.5)
52:54:00:a3:46:7f (131.252.208.60)
52:54:00:a9:30:9f (131.252.208.117)
52:54:00:a9:30:9f (131.252.208.53)
52:54:00:b4:6e:05 (131.252.208.34)
52:54:00:c2:05:63 (131.252.208.137)
52:54:00:cf:4c:1b (131.252.208.36)
52:54:00:eb:9a:42 (131.252.208.28)
52:54:00:f2:09:bc (131.252.208.85)
52:54:00:f6:f8:54 (131.252.208.54)
cc:aa:77:07:f2:7a (131.252.208.171)
cc:aa:77:0b:76:be (131.252.208.77)
cc:aa:77:50:b9:5d (131.252.208.17)
cc:aa:77:5f:de:0e (131.252.208.110)
cc:aa:77:ed:72:3e (131.252.208.43)
cc:aa:77:f1:d3:21 (131.252.208.63)
agah@ada:~$ arp -a | sort -k 4 | awk '{print $4, $2}' | uniq | wc -l
25
agah@ada:~$

```

- create a file that contains each IP address that appears in the machine's ARP table and places the results in a file called `arp_entries`

```
arp -an | awk -F '[]' '{print $2}' > arp_entries
```

- What network prefix do most of the IP addresses in the ARP table share?

In the `arp_entries` output, most of the IP addresses share the network prefix `131.252.208`. This prefix represents the first three octets of each IP address in the list, indicating that these addresses are likely on the same subnet within the `131.252.208.0/24` range.

```
agah@ada:~$ cat arp_entries
131.252.208.34
131.252.208.20
131.252.208.63
131.252.208.43
131.252.208.23
131.252.208.17
131.252.208.117
131.252.208.138
131.252.208.83
131.252.208.36
131.252.208.28
131.252.208.94
131.252.208.53
131.252.208.5
131.252.208.11
131.252.208.85
131.252.208.77
131.252.208.171
131.252.208.137
131.252.208.84
131.252.208.110
131.252.208.1
131.252.208.81
131.252.208.54
131.252.208.60
agah@ada:~$
```

1.2.3 ARP (Cloud)

- Find the IP address and hardware address of the local ethernet card interface (Typically beginning with `eth`, `ens`, or `enp`).

```

ahmadagah@course-vm:~/Desktop/cloud-agah-agah$ ip address
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 42:01:0a:8a:00:02 brd ff:ff:ff:ff:ff:ff
    inet 10.138.0.2/32 metric 100 scope global dynamic ens4
        valid_lft 82096sec preferred_lft 82096sec
    inet6 fe80::4001:aff:fe8a:2/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:b0:f6:25:6f brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
ahmadagah@course-vm:~/Desktop/cloud-agah-agah$

```

- What is the default router's IP address (e.g. the gateway address for the default route 0.0.0.0/0)

```

ahmadagah@course-vm:~/Desktop/cloud-agah-agah$ netstat -rn
Kernel IP routing table
Destination    Gateway         Genmask         Flags   MSS Window  irtt Iface
0.0.0.0        10.138.0.1     0.0.0.0         UG          0  0          0 ens4
10.138.0.1     0.0.0.0        255.255.255.255 UH          0  0          0 ens4
169.254.169.254 10.138.0.1    255.255.255.255 UGH         0  0          0 ens4
172.17.0.0     0.0.0.0        255.255.0.0     U          0  0          0 docker0
ahmadagah@course-vm:~/Desktop/cloud-agah-agah$

```

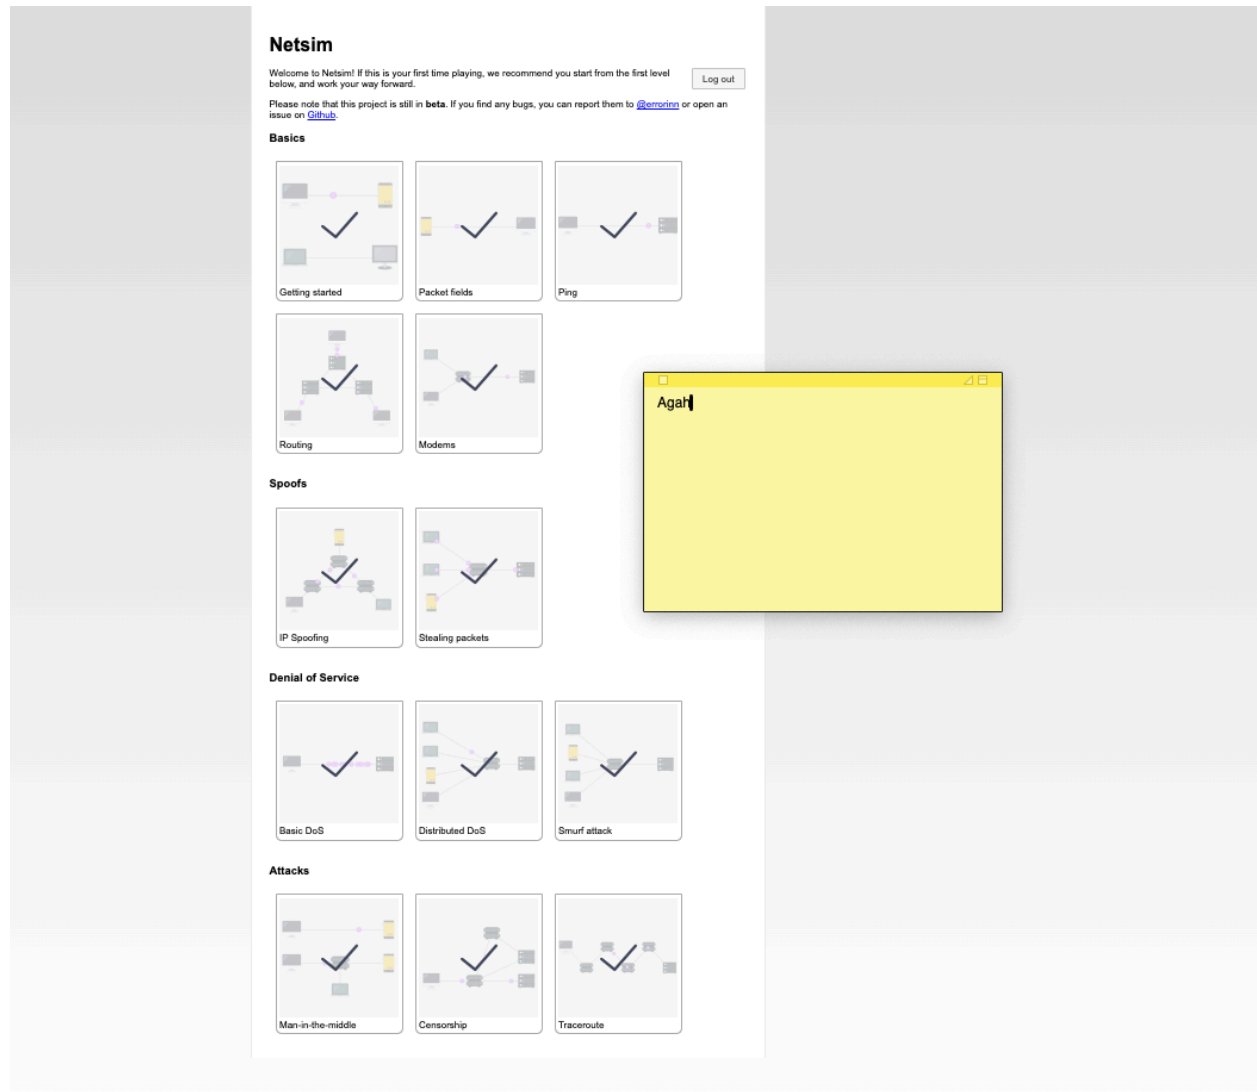
- What is the default router's hardware address?

```

ahmadagah@course-vm:~/Desktop/cloud-agah-agah$ arp 10.138.0.1
Address          HWtype  HWaddress      Flags Mask    Iface
_gateway         ether   42:01:0a:8a:00:01 C              ens4

```

1.2.4 Netsim



1.3 Cloud Computing

1.3.3 Scan targets for services

- run `nmap` on the internal subnet the instances have been placed on:

```

ahmadagah@course-vm:~/Desktop/cloud-agah-agah$ sudo nmap 10.138.0.0/20
Starting Nmap 7.80 ( https://nmap.org ) at 2024-10-08 06:03 UTC
Nmap scan report for apache-1-vm.c.cloud-agah-agah.internal (10.138.0.3)
Host is up (0.00010s latency).
Not shown: 998 closed ports
PORT      STATE SERVICE
22/tcp    open  ssh
80/tcp    open  http

Nmap scan report for django-1-vm.c.cloud-agah-agah.internal (10.138.0.4)
Host is up (0.000078s latency).
Not shown: 999 closed ports
PORT      STATE SERVICE
22/tcp    open  ssh

Nmap scan report for secured-wordpress-on-ubuntu-14-04-lts-1-vm.c.cloud-agah-agah.internal (10.138.0.5)
Host is up (0.000092s latency).
Not shown: 998 closed ports
PORT      STATE SERVICE
22/tcp    open  ssh
80/tcp    open  http

Nmap scan report for _gateway (10.138.0.1)
Host is up (0.0011s latency).
Not shown: 999 filtered ports
PORT      STATE SERVICE
53/tcp    open  domain
MAC Address: 42:01:0A:8A:00:01 (Unknown)

Nmap scan report for course-vm.c.cloud-agah-agah.internal (10.138.0.2)
Host is up (0.0000080s latency).
Not shown: 998 closed ports
PORT      STATE SERVICE
22/tcp    open  ssh
3389/tcp  open  ms-wbt-server

Nmap done: 4096 IP addresses (5 hosts up) scanned in 45.35 seconds
ahmadagah@course-vm:~/Desktop/cloud-agah-agah$

```

1.3.5 Navigating default networks

- How many subnetworks are created initially on the `default` network? How many regions does this correspond to? (Use a pipe to pass output to `grep` in order to return specific lines of output and then another to pass output to `wc` to count them: `| grep default | wc -l`)


```

~ gcloud compute networks subnets list
NAME          REGION          NETWORK RANGE      STACK_TYPE  IPV6_ACCESS_TYPE  INTERNAL_IPV6_PREFIX  EXTERNAL_IPV6_PREFIX
default       us-central1     default 10.128.0.0/20  IPV4_ONLY
default       europe-west1    default 10.132.0.0/20  IPV4_ONLY
default       us-west1        default 10.138.0.0/20  IPV4_ONLY
default       asia-east1      default 10.140.0.0/20  IPV4_ONLY
default       us-east1        default 10.142.0.0/20  IPV4_ONLY
default       asia-northeast1 default 10.146.0.0/20  IPV4_ONLY
default       asia-southeast1 default 10.148.0.0/20  IPV4_ONLY
default       us-east4        default 10.150.0.0/20  IPV4_ONLY
default       australia-southeast1 default 10.152.0.0/20  IPV4_ONLY
default       europe-west2    default 10.154.0.0/20  IPV4_ONLY
default       europe-west3    default 10.156.0.0/20  IPV4_ONLY
default       southamerica-east1 default 10.158.0.0/20  IPV4_ONLY
default       asia-south1     default 10.160.0.0/20  IPV4_ONLY
default       northamerica-northeast1 default 10.162.0.0/20  IPV4_ONLY
default       europe-west4    default 10.164.0.0/20  IPV4_ONLY
default       europe-north1   default 10.166.0.0/20  IPV4_ONLY
default       us-west2        default 10.168.0.0/20  IPV4_ONLY
default       asia-east2      default 10.170.0.0/20  IPV4_ONLY
default       europe-west6    default 10.172.0.0/20  IPV4_ONLY
default       asia-northeast2 default 10.174.0.0/20  IPV4_ONLY
default       asia-northeast3 default 10.178.0.0/20  IPV4_ONLY
default       us-west3        default 10.180.0.0/20  IPV4_ONLY
default       us-west4        default 10.182.0.0/20  IPV4_ONLY
default       asia-southeast2 default 10.184.0.0/20  IPV4_ONLY
default       europe-central2 default 10.186.0.0/20  IPV4_ONLY
default       northamerica-northeast2 default 10.188.0.0/20  IPV4_ONLY
default       asia-south2     default 10.190.0.0/20  IPV4_ONLY
default       australia-southeast2 default 10.192.0.0/20  IPV4_ONLY
default       southamerica-west1 default 10.194.0.0/20  IPV4_ONLY
default       us-east7        default 10.196.0.0/20  IPV4_ONLY
default       europe-west8    default 10.198.0.0/20  IPV4_ONLY
default       europe-west9    default 10.200.0.0/20  IPV4_ONLY
default       us-east5        default 10.202.0.0/20  IPV4_ONLY
default       europe-southwest1 default 10.204.0.0/20  IPV4_ONLY
default       us-south1       default 10.206.0.0/20  IPV4_ONLY
default       me-west1        default 10.208.0.0/20  IPV4_ONLY
default       europe-west12   default 10.210.0.0/20  IPV4_ONLY
default       me-central1     default 10.212.0.0/20  IPV4_ONLY
default       europe-west10   default 10.214.0.0/20  IPV4_ONLY
default       africa-south1   default 10.218.0.0/20  IPV4_ONLY
default       us-west8        default 10.220.0.0/20  IPV4_ONLY
default       northamerica-south1 default 10.224.0.0/20  IPV4_ONLY
~ gcloud compute networks subnets list | grep default | wc -l
42
~

```

there are **42 subnetworks** created initially on the **default network**. Each line with **default** represents a different **subnetwork**.

Since each subnetwork is in a specific **region**, the number of subnetworks corresponds directly to the number of regions for the **default network**.

- Given the CIDR prefix associated with each subnetwork, how many hosts does each subnetwork support?
- Total Bits for IP Addressing:** IPv4 addresses have 32 bits.
- Network Bits in /20:** The **/20** prefix means the first 20 bits are used for the network.
- Host Bits:** $32 - 20 = 12$ bits are available for host addresses.
- Calculating Hosts:**

- The number of host addresses is calculated by $2^{12} - 2 = 4096 - 2 = 4094$.
- 2 addresses are subtracted because one address is reserved for the network address and one for the broadcast address.
- Therefore, $2^{12} - 2 = 4096 - 2 = 4094$.

- Create two instances in [different zones in separate regions](#) of your choice:
- List both instances.
- Which CIDR subnetworks are these instances brought up in? Do they correspond to the appropriate region based on the prior commands?

```
➔ ~ gcloud compute instances list
```

NAME	ZONE	MACHINE_TYPE	PREEMPTIBLE	INTERNAL_IP	EXTERNAL_IP	STATUS
course-vm	us-west1-b	e2-medium		10.138.0.2	34.19.41.230	RUNNING
instance-2	us-east1-b	n1-standard-1		10.142.0.2	35.237.127.195	RUNNING
instance-1	us-east4-c	n1-standard-1		10.150.0.2	34.21.32.117	RUNNING

Instance: **course-vm**

- **Zone:** **us-west1-b**
- **CIDR Subnetwork:** **10.138.0.0/20**
- **Corresponding Region:** **us-west1**
- **Check:** This matches the subnet range **10.138.0.0/20** for the **us-west1** region from the subnetwork list.

Instance: **instance-2**

- **Zone:** **us-east1-b**
- **CIDR Subnetwork:** **10.142.0.0/20**
- **Corresponding Region:** **us-east1**
- **Check:** This matches the subnet range **10.142.0.0/20** for the **us-east1** region from the subnetwork list.

Instance: **instance-1**

- **Zone:** **us-east4-c**
- **CIDR Subnetwork:** **10.150.0.0/20**
- **Corresponding Region:** **us-east4**
- **Check:** This matches the subnet range **10.150.0.0/20** for the **us-east4** region from the subnetwork list.

- From `instance-1`, perform a `ping` to the `Internal IP` address of `instance-2`. Take a screenshot of the output.
- From the figure in the previous step. What facilitates this connectivity: the virtual switch or the VPN Gateway?

```
ahmadagah@instance-1:~$ ping -c 5 10.142.0.2
PING 10.142.0.2 (10.142.0.2) 56(84) bytes of data.
64 bytes from 10.142.0.2: icmp_seq=1 ttl=64 time=13.8 ms
64 bytes from 10.142.0.2: icmp_seq=2 ttl=64 time=12.8 ms
64 bytes from 10.142.0.2: icmp_seq=3 ttl=64 time=13.2 ms
64 bytes from 10.142.0.2: icmp_seq=4 ttl=64 time=12.8 ms
64 bytes from 10.142.0.2: icmp_seq=5 ttl=64 time=12.8 ms

--- 10.142.0.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4005ms
rtt min/avg/max/mdev = 12.798/13.101/13.831/0.392 ms
ahmadagah@instance-1:~$
```

The **virtual switch** within the Google Cloud VPC is what facilitates the internal connectivity shown in `ping` command, allowing direct communication between `10.142.0.2` and other instances within the same VPC network.

1.3.6 Creating custom networks

- Take a screenshot of the new subnets created in `custom-network1` alongside the default subnetworks in those regions assigned to the `default` network.

```
+ ~ gcloud compute networks subnets list --regions=us-central1,europe-west1
NAME          REGION    NETWORK    RANGE          STACK_TYPE  IPV6_ACCESS_TYPE  INTERNAL_IPV6_PREFIX  EXTERNAL_IPV6_PREFIX
default       europe-west1  default    10.132.0.0/20  IPV4_ONLY
subnet-europe-west-192  europe-west1  custom-network1  192.168.5.0/24  IPV4_ONLY
default       us-central1  default    10.128.0.0/20  IPV4_ONLY
subnet-us-central-192  us-central1  custom-network1  192.168.1.0/24  IPV4_ONLY
```

- Explain why the result of this ping is different from when you performed the ping to `instance-2`.

```

ahmadagah@course-vm:~$ ping -c 5 192.168.1.2
PING 192.168.1.2 (192.168.1.2) 56(84) bytes of data.

--- 192.168.1.2 ping statistics ---
5 packets transmitted, 0 received, 100% packet loss, time 4089ms

ahmadagah@course-vm:~$ ping -c 5 192.168.5.2
PING 192.168.5.2 (192.168.5.2) 56(84) bytes of data.

--- 192.168.5.2 ping statistics ---
5 packets transmitted, 0 received, 100% packet loss, time 4100ms

```

The successful **ping** to **instance-2** (10.142.0.2) occurred within the **10.x.y.z** IP range, which is part of the **default VPC network** in Google Cloud. Instances in the same VPC network, Google Cloud's default or custom subnetworks, can communicate with each other over their internal IP addresses.

The IP addresses **192.168.1.2** and **192.168.5.2** belong to a **different IP range** that is not part of the **default Google Cloud VPC network**. These are in a private IP range commonly used for internal networks, but they are not accessible from the **10.x.y.z** subnet without additional routing or VPN configurations.

- Take screenshots of all 4 instances in the UI including the network they belong to.

<input type="checkbox"/>	Status	Name ↑	Zone	Recommendations	In use by	Internal IP	External IP	Network	Connect
<input type="checkbox"/>	✓	course-vm	us-west1-b			10.138.0.2 (nic0)	34.19.41.230 (nic0)	default	SSH ▾ ⋮
<input type="checkbox"/>	✓	instance-1	us-east4-c			10.150.0.2 (nic0)	34.21.32.117 (nic0)	default	SSH ▾ ⋮
<input type="checkbox"/>	✓	instance-2	us-east1-b			10.142.0.2 (nic0)	35.237.127.195 (nic0)	default	SSH ▾ ⋮
<input type="checkbox"/>	✓	instance-3	us-central1-a			192.168.1.2 (nic0)	34.122.145.160 (nic0)	custom-network1	SSH ▾ ⋮
<input type="checkbox"/>	✓	instance-4	europa-west1-d			192.168.5.2 (nic0)	34.78.11.54 (nic0)	custom-network1	SSH ▾ ⋮

- Take a screenshot of the subnetworks created for the **custom-network1** network and some of the subnetworks of the **default** network showing their regions, internal IP ranges and Gateways.

default

OVERVIEW SUBNETS STATIC INTERNAL IP ADDRESSES FIREWALLS FIREWALL ENDPOINTS ROUTES VPC NETWORK PEERING PRIVATE SERVICES ACCESS DNS CONFIGURATION

Subnets

ADD SUBNET MANAGE FLOW LOGS

Filter	Enter property name or value											
	Name	Region	Stack Type	Primary IPv4 range	Secondary IPv4 ranges	IPv6 ranges	Reserved internal ranges	Gateway	Private Google Access	Flow logs		
	default	afrika-south1	IPv4	10.218.0.0/20			None	10.218.0.1	Off	Off		
	default	asia-east1	IPv4	10.140.0.0/20			None	10.140.0.1	Off	Off		
	default	asia-east2	IPv4	10.170.0.0/20			None	10.170.0.1	Off	Off		
	default	asia-northeast1	IPv4	10.146.0.0/20			None	10.146.0.1	Off	Off		
	default	asia-northeast2	IPv4	10.174.0.0/20			None	10.174.0.1	Off	Off		
	default	asia-northeast3	IPv4	10.178.0.0/20			None	10.178.0.1	Off	Off		
	default	asia-south1	IPv4	10.160.0.0/20			None	10.160.0.1	Off	Off		
	default	asia-south2	IPv4	10.190.0.0/20			None	10.190.0.1	Off	Off		
	default	asia-southeast1	IPv4	10.148.0.0/20			None	10.148.0.1	Off	Off		
	default	asia-southeast2	IPv4	10.184.0.0/20			None	10.184.0.1	Off	Off		
	default	australia-southeast1	IPv4	10.152.0.0/20			None	10.152.0.1	Off	Off		
	default	australia-southeast2	IPv4	10.192.0.0/20			None	10.192.0.1	Off	Off		
	default	eu-central1	IPv4	10.186.0.0/20			None	10.186.0.1	Off	Off		
	default	eu-north1	IPv4	10.166.0.0/20			None	10.166.0.1	Off	Off		
	default	eu-southwest1	IPv4	10.204.0.0/20			None	10.204.0.1	Off	Off		
	default	eu-west1	IPv4	10.132.0.0/20			None	10.132.0.1	Off	Off		
	default	eu-west10	IPv4	10.214.0.0/20			None	10.214.0.1	Off	Off		
	default	eu-west12	IPv4	10.210.0.0/20			None	10.210.0.1	Off	Off		
	default	eu-west2	IPv4	10.154.0.0/20			None	10.154.0.1	Off	Off		
	default	eu-west3	IPv4	10.156.0.0/20			None	10.156.0.1	Off	Off		
	default	eu-west4	IPv4	10.164.0.0/20			None	10.164.0.1	Off	Off		
	default	eu-west6	IPv4	10.172.0.0/20			None	10.172.0.1	Off	Off		
	default	eu-west8	IPv4	10.198.0.0/20			None	10.198.0.1	Off	Off		
	default	eu-west9	IPv4	10.200.0.0/20			None	10.200.0.1	Off	Off		
	default	eu-west11	IPv4	10.255.0.0/20			None	10.255.0.1	Off	Off		

custom-network1

OVERVIEW SUBNETS STATIC INTERNAL IP ADDRESSES FIREWALLS FIREWALL ENDPOINTS ROUTES VPC NETWORK PEERING PRIVATE SERVICES ACCESS DNS CONFIGURATION

Subnets

ADD SUBNET MANAGE FLOW LOGS

Filter	Enter property name or value											
	Name	Region	Stack Type	Primary IPv4 range	Secondary IPv4 ranges	IPv6 ranges	Reserved internal ranges	Gateway	Private Google Access	Flow logs		
	subnet-europe-west-192	eu-central1	IPv4	192.168.5.0/24			None	192.168.5.1	Off	Off		
	subnet-us-central-192	us-central1	IPv4	192.168.1.0/24			None	192.168.1.1	Off	Off		

Reserved proxy-only subnets for load balancing

